

F1_analysis

2025-12-07

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com> (<http://rmarkdown.rstudio.com>).

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
#Shiv
#setwd("/Users/shivpatel/Desktop/Code")

# -----
# Libraries
# -----
library(httr)
library(jsonlite)
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##     filter, lag
```

```
## The following objects are masked from 'package:base':
##     intersect, setdiff, setequal, union
```

```
library(purrr)
```

```
##
## Attaching package: 'purrr'
```

```
## The following object is masked from 'package:jsonlite':
##     flatten
```

```
library(tidyr)
library(openmeteo)
library(lubridate)
```

```
##
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union
```

```
library(factoextra)
```

```
## Loading required package: ggplot2
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WB
```

```
library(rpart)
library(partykit)
```

```
## Loading required package: grid
```

```
## Loading required package: libcoin
```

```
## Loading required package: mvtnorm
```

```
library(party)
```

```
## Loading required package: modeltools
```

```
## Loading required package: stats4
```

```
## Loading required package: strucchange
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##     as.Date, as.Date.numeric
```

```
## Loading required package: sandwich
```

```
##  
## Attaching package: 'party'
```

```
## The following objects are masked from 'package:partykit':  
##  
##     cforest, ctree, ctree_control, edge_simple, mob, mob_control,  
##     node_barplot, node_bivplot, node_boxplot, node_inner, node_surv,  
##     node_terminal, varimp
```

```
## The following object is masked from 'package:dplyr':  
##  
##     where
```

```
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':  
##   method      from  
##   as.zoo.data.frame zoo
```

```

# -----
# Function to import a given years data frame and then join it onto master
# -----
import_clean <- function(frames,year){
  # -----
  # Get all Race sessions for the year
  # -----
  sessions_url <- paste0("https://api.openf1.org/v1/sessions?year=",year,"&session_name=Race")

  sessions_raw <- fromJSON(content(GET(sessions_url), "text", encoding = "UTF-8"))
  sessions <- sessions_raw

  # -----
  # Position and Driver Data
  # -----
  results <- map_df(sessions_raw$session_key,
                     function(ses){
    Sys.sleep(0.5)
    url <- paste0('https://api.openf1.org/v1/drivers?session_key=',ses)
    df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
    df$session_key <- ses
    return(df)})

  drivers <- results

  results <- map_df(sessions_raw$session_key,
                     function(ses){
    Sys.sleep(0.5)
    url <- paste0("https://api.openf1.org/v1/position?session_key=",ses)
    df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
    df$session_key <- ses
    return(df)})

  position <- results

  driver_pos <- left_join(position, drivers, by = join_by(session_key==session_key, driver_number==driver_number))
  driver_pos <- select(driver_pos,"date","session_key","driver_number","position","full_name","team_name")

  driver_pos <- driver_pos %>%
    group_by(session_key,driver_number) %>%
    filter(date == max(date)) %>%
    pivot_wider(
      id_cols = session_key,
      names_from = position,
      values_from = c(driver_number,full_name,team_name),
      values_fn = list)
}

```

```

names_prefix = "pos_"

# -----
# Weather
# -----
results <- map_df(sessions_raw$session_key,
                  function(ses){
                    Sys.sleep(0.5)
                    url <- paste0('https://api.openf1.org/v1/weather?session_key=' , se
s)
                    df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
                    df$session_key <- ses
                    return(df)})}

weather <- results

weather <- weather %>%
  group_by(session_key) %>%
  summarize(wind_direction = mean(wind_direction, na.rm = TRUE),
            meeting_key = max(meeting_key, na.rm = TRUE),
            wind_speed = mean(wind_speed, na.rm = TRUE),
            #Rainfall is normally a binary denoteing if in a given minute there was r
ain. This new rainfall variable demotes the percentage of minutes in the race that had r
ain
            track_temperature = mean(track_temperature, na.rm = TRUE),
            rainfall = sum(rainfall, na.rm = TRUE)/length(rainfall),
            air_temperature = mean(air_temperature, na.rm = TRUE),
            humidity = mean(humidity, na.rm = TRUE),
            pressure = mean(pressure, na.rm = TRUE))

# -----
# Get all qualifying sessions for the year Qualifying
# -----
results <- map_df(sessions_raw$meeting_key,
                  function(meet){
                    Sys.sleep(0.5)
                    url <- paste0("https://api.openf1.org/v1/sessions?&session_name=Qu
alifying&meeting_key=",meet)
                    df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
                    df$meeting_key <- meet
                    return(df)})}

quals <- results[c(1,2)]

# -----
# Qual_results
# -----
results <- map_df(quals$session_key,
                  function(ses){
                    Sys.sleep(0.5)

```

```

        url <- paste0("https://api.openf1.org/v1/position?session_key=", se
s)
        df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
        #df$meeting_key <- ses
        return(df)})}

quals_results <- results

quals_results <- left_join(quals_results, drivers, by = join_by(meeting_key==meeting_
key, driver_number==driver_number)) %>%
  select("date","meeting_key","session_key.x","driver_number","position","full_nam
e","team_name")

quals_results <- quals_results %>%
  group_by(session_key.x,driver_number) %>%
  filter(date == max(date))
quals_results_2 <- quals_results
quals_results <- pivot_wider(quals_results,
  id_cols = c(session_key.x,meeting_key),
  names_from = position,
  values_from = c(driver_number,full_name,team_name),
  names_prefix = "qual_pos_")

# -----
# 4. Merge into final DF
# -----

full_df <- sessions %>%
  left_join(driver_pos, by = "session_key") %>%
  left_join(weather, by = "session_key")%>%
  left_join(quals_results, by = c("meeting_key.x"="meeting_key"))

full_df <- select(full_df,meeting_key = "meeting_key.x",
  "session_key.x",
  "location",
  "date_start",
  "date_end",
  "country_key",
  "country_name",
  "circuit_key",
  "circuit_short_name",
  "year",
  "driver_number_pos_20",
  "driver_number_pos_19",
  "driver_number_pos_18",
  "driver_number_pos_17",
  "driver_number_pos_16",
  "driver_number_pos_15",
  "driver_number_pos_14",
  "driver_number_pos_13",
  "driver_number_pos_12",
  "driver_number_pos_11",
  "driver_number_pos_10",
  "driver_number_pos_9",
  "driver_number_pos_8",
  "driver_number_pos_7",
  "driver_number_pos_6",
  "driver_number_pos_5",
  "driver_number_pos_4",
  "driver_number_pos_3",
  "driver_number_pos_2",
  "driver_number_pos_1")

```

"driver_number_pos_10",
"driver_number_pos_9",
"driver_number_pos_8",
"driver_number_pos_7",
"driver_number_pos_6",
"driver_number_pos_5",
"driver_number_pos_4",
"driver_number_pos_3",
"driver_number_pos_2",
"driver_number_pos_1",
"full_name_pos_20",
"full_name_pos_19",
"full_name_pos_18",
"full_name_pos_17",
"full_name_pos_16",
"full_name_pos_15",
"full_name_pos_14",
"full_name_pos_13",
"full_name_pos_12",
"full_name_pos_11",
"full_name_pos_10",
"full_name_pos_9",
"full_name_pos_8",
"full_name_pos_7",
"full_name_pos_6",
"full_name_pos_5",
"full_name_pos_4",
"full_name_pos_3",
"full_name_pos_2",
"full_name_pos_1",
"team_name_pos_20",
"team_name_pos_19",
"team_name_pos_18",
"team_name_pos_17",
"team_name_pos_16",
"team_name_pos_15",
"team_name_pos_14",
"team_name_pos_13",
"team_name_pos_12",
"team_name_pos_11",
"team_name_pos_10",
"team_name_pos_9" ,
"team_name_pos_8",
"team_name_pos_7",
"team_name_pos_6" ,
"team_name_pos_5",
"team_name_pos_4",
"team_name_pos_3",
"team_name_pos_2",
"team_name_pos_1",
"driver_number_qual_pos_20",
"driver_number_qual_pos_19",

"driver_number_qual_pos_18",
"driver_number_qual_pos_17",
"driver_number_qual_pos_16",
"driver_number_qual_pos_15",
"driver_number_qual_pos_14",
"driver_number_qual_pos_13",
"driver_number_qual_pos_12",
"driver_number_qual_pos_11",
"driver_number_qual_pos_10",
"driver_number_qual_pos_9",
"driver_number_qual_pos_8",
"driver_number_qual_pos_7",
"driver_number_qual_pos_6",
"driver_number_qual_pos_5",
"driver_number_qual_pos_4",
"driver_number_qual_pos_3",
"driver_number_qual_pos_2",
"driver_number_qual_pos_1",
"full_name_qual_pos_20",
"full_name_qual_pos_19",
"full_name_qual_pos_18",
"full_name_qual_pos_17",
"full_name_qual_pos_16",
"full_name_qual_pos_15",
"full_name_qual_pos_14",
"full_name_qual_pos_13",
"full_name_qual_pos_12",
"full_name_qual_pos_11",
"full_name_qual_pos_10",
"full_name_qual_pos_9",
"full_name_qual_pos_8",
"full_name_qual_pos_7",
"full_name_qual_pos_6",
"full_name_qual_pos_5",
"full_name_qual_pos_4",
"full_name_qual_pos_3",
"full_name_qual_pos_2",
"full_name_qual_pos_1",
"team_name_qual_pos_20",
"team_name_qual_pos_19",
"team_name_qual_pos_18",
"team_name_qual_pos_17",
"team_name_qual_pos_16",
"team_name_qual_pos_15",
"team_name_qual_pos_14",
"team_name_qual_pos_13",
"team_name_qual_pos_12",
"team_name_qual_pos_11",
"team_name_qual_pos_10",
"team_name_qual_pos_9",
"team_name_qual_pos_8",
"team_name_qual_pos_7",

```

        "team_name_qual_pos_6" ,
        "team_name_qual_pos_5",
        "team_name_qual_pos_4",
        "team_name_qual_pos_3",
        "team_name_qual_pos_2",
        "team_name_qual_pos_1",
        "wind_direction",
        "wind_speed",
        "rainfall",
        "track_temperature",
        "air_temperature",
        "humidity",
        "pressure")

master_df_tall <- left_join(drivers,position,join_by(session_key,meeting_key,driver_number)) %>%
  left_join(sessions,join_by(session_key,meeting_key)) %>%
  left_join(weather,join_by(session_key,meeting_key)) %>%
  left_join(quals_results_2,join_by(meeting_key,driver_number))

return(list(rbind(as.data.frame(frames[1]),full_df),rbind(as.data.frame(frames[2]),master_df_tall)))}

# -----
# Import our needed years
# -----
## Import if not saved locally
frames <- list(data.frame(),data.frame())
frames <- import_clean(frames,2023)
frames <- import_clean(frames,2024)
frames <- import_clean(frames,2025)

master_df <- as.data.frame(frames[1])
master_df_tall <- as.data.frame(frames[2])
#
# saveRDS(master_df, "master_hist_table.rds")
# saveRDS(master_df_tall, "master_hist_table_tall.rds")

#Import locally stored file (saves API calls)
#master_df <- readRDS("master_hist_table.rds")
#master_df_tall <- readRDS("master_hist_table_tall.rds")
#Shiv

# -----
# Clustering locations
# -----
#Temperature difference may be more useful than individual temps
master_df$temp_diff <- master_df$track_temperature - master_df$air_temperature

#Aggregate track temps
track_features_agg <- master_df %>%
  group_by(location) %>%

```

```

summarise(
  Avg_Wind_Speed = mean(wind_speed, na.rm = TRUE),
  Avg_Rainfall_Perc = mean(rainfall, na.rm = TRUE),
  Avg_Track_Temp = mean(track_temperature, na.rm = TRUE),
  Avg_Air_Temp = mean(air_temperature, na.rm = TRUE),
  Avg_Humidity = mean(humidity, na.rm = TRUE),
  Avg_Pressure = mean(pressure, na.rm = TRUE),
  Avg_Temp_Diff = mean(temp_diff, na.rm = TRUE),
  .groups = 'drop'
)

clustering_data <- track_features_agg %>%
  select(starts_with("Avg_")) %>%
  na.omit()

scaled_data <- scale(clustering_data)

wss <- numeric(20)
for (k in 1:20) {
  kmeans_result <- kmeans(scaled_data, centers = k, nstart = 25, iter.max = 30)
  wss[k] <- kmeans_result$tot.withinss
}
print("WSS Values for k=1 to 10:")

## [1] "WSS Values for k=1 to 10:"
```

```

print(wss)
```


```

## [1] 161.00000 114.542128 85.036998 66.785549 52.738398 44.378923
## [7] 36.914603 29.538929 23.957175 20.145769 16.906088 14.248345
## [13] 11.962048 10.158705 8.612900 7.210594 5.852706 4.718458
## [19] 3.684299 2.822952
```

```


plot(1:20, wss, type="b", xlab="Number of Clusters (k)", ylab="Within Sum of Squares (WS)", main="Elbow Method for Optimal k")
```

Elbow Method for Optimal k



```
fviz_nbclust(scaled_data, FUNcluster = kmeans)
```

Optimal number of clusters



```
k_optimal <- 2

final_clusters <- kmeans(scaled_data, centers = k_optimal, nstart = 25, iter.max = 30)

cluster_assignments <- track_features_agg %>%
  filter(complete.cases(select(., starts_with("Avg_")))) %>%
  mutate(Track_Cluster = final_clusters$cluster) %>%
  select(location, Track_Cluster, everything())

print(cluster_assignments)
```

```

## # A tibble: 24 × 9
##   location  Track_Cluster Avg_Wind_Speed Avg_Rainfall_Perc Avg_Track_Temp
##   <chr>          <int>        <dbl>            <dbl>            <dbl>
## 1 Austin           1         2.30             0              44.1
## 2 Baku             1         1.38             0              37.2
## 3 Barcelona         1         1.88            0.00216          41.2
## 4 Budapest          1         1.59             0              41.8
## 5 Imola            1         2.69             0              42.3
## 6 Jeddah            2         1.37             0              33.7
## 7 Las Vegas         2         2.21             0              18.4
## 8 Lusail            2         1.27             0              29.3
## 9 Marina Bay        2         1.16            0.0124          35.6
## 10 Melbourne         2         1.85            0.109           29.2
## # i 14 more rows
## # i 4 more variables: Avg_Air_Temp <dbl>, Avg_Humidity <dbl>,
## #   Avg_Pressure <dbl>, Avg_Temp_Diff <dbl>

```


```

pca <- prcomp(scaled_data)
pca_df <- data.frame(
  PC1 = pca$x[, 1],
  PC2 = pca$x[, 2],
  Cluster = cluster_assignments$Track_Cluster
)

ggplot(pca_df, aes(PC1, PC2, color = Cluster)) +
  geom_point(size = 3) +
  theme_minimal() +
  ggtitle("K-Means Clusters (PCA Projection)")

```

K-Means Clusters (PCA Projection)



```

master_df <- left_join(master_df,cluster_assignments,by="location")

# -----
# Add in additional variables
# -----
analysis_set <- left_join(master_df_tall,cluster_assignments,by="location")
analysis_set <- select(analysis_set,"driver_number",
                      full_name="full_name.x",
                      team_name = "team_name.x",
                      race_position = "position.x",
                      "location",
                      "year",
                      "wind_speed",
                      "track_temperature",
                      "rainfall",
                      "air_temperature",
                      "humidity",
                      "pressure",
                      qual_position = "position.y",
                      loc_cluster = Track_Cluster)
analysis_set$temp_diff <- analysis_set$track_temperature - analysis_set$air_temperature
analysis_set$driver_number <- as.factor(analysis_set$driver_number)
analysis_set$full_name <- as.factor(analysis_set$full_name)
analysis_set$team_name <- as.factor(analysis_set$team_name)
analysis_set$location <- as.factor(analysis_set$location)
analysis_set$year <- as.factor(analysis_set$year)
analysis_set$loc_cluster <- as.factor(analysis_set$loc_cluster)


# -----
# Make Train and Test Sets
# -----
set.seed(112224)
ind <- sample(1:nrow(analysis_set), size = 0.7*nrow(analysis_set), replace = FALSE)
test <- analysis_set[-ind,]
train <- analysis_set[ind,]

# -----
# Linear Model
# -----
# ##

# ##
#First look for inter-relationship patterns

```

```
# ##
pairs(~., data=train[-c(1,6,8,10)], pch=".")
```



```
# ##
# Model using qualifying results
# ##

lm_wqual <- lm(race_position ~ ., data = train[c("team_name","wind_speed","rainfall","humidity","pressure","qual_position","loc_cluster","temp_diff","race_position")])
#N:Obviously large amount of multicolinearity between name and team
summary(lm_wqual)
```

```

## 
## Call:
## lm(formula = race_position ~ ., data = train[c("team_name", "wind_speed",
## "rainfall", "humidity", "pressure", "qual_position", "loc_cluster",
## "temp_diff", "race_position")])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.5659  -2.4499  -0.2054   2.2924  14.4609
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)             5.6973695  0.4846213 11.756 < 2e-16 ***
## team_nameAlphaTauri     -0.4602610  0.1620438 -2.840 0.004510 **  
## team_nameAlpine          -0.3645462  0.1315489 -2.771 0.005590 **  
## team_nameAston Martin    -0.7557428  0.1318753 -5.731 1.01e-08 ***  
## team_nameFerrari         -2.4410510  0.1438583 -16.968 < 2e-16 ***  
## team_nameHaas F1 Team    -0.2399908  0.1296182 -1.852 0.064107 .    
## team_nameKick Sauber      0.3745841  0.1422188  2.634 0.008448 **  
## team_nameMcLaren          -2.0402580  0.1438182 -14.186 < 2e-16 ***  
## team_nameMercedes         -1.7280713  0.1396832 -12.371 < 2e-16 ***  
## team_nameRacing Bulls     0.1275074  0.1717690  0.742 0.457902    
## team_nameRB                0.1649315  0.1653202  0.998 0.318460    
## team_nameRed Bull Racing   -2.3387322  0.1403298 -16.666 < 2e-16 ***  
## team_nameWilliams         -0.1639976  0.1315355 -1.247 0.212485    
## wind_speed                  0.0113492  0.0301699  0.376 0.706789    
## rainfall                     0.2097703  0.2713171  0.773 0.439438    
## humidity                    -0.0023706  0.0018234 -1.300 0.193573    
## pressure                     0.0018329  0.0004721  3.882 0.000104 ***  
## qual_position                 0.4557103  0.0052902  86.142 < 2e-16 ***  
## loc_cluster2                 -0.2660195  0.0592456 -4.490 7.15e-06 ***  
## temp_diff                     -0.0186156  0.0045725 -4.071 4.69e-05 ***  
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.538 on 23574 degrees of freedom
## Multiple R-squared:  0.4347, Adjusted R-squared:  0.4342 
## F-statistic:  954 on 19 and 23574 DF,  p-value: < 2.2e-16

```

#Step

```

lm_race_step_wqual <- step(lm_wqual, scope=list(lower=lm(train$race_position~0), upper=lm_wqual), direction = "forward")

```

```

## Start: AIC=59639.09
## race_position ~ team_name + wind_speed + rainfall + humidity +
##               pressure + qual_position + loc_cluster + temp_diff

```

```

summary(lm_race_step_wqual)

```

```

## 
## Call:
## lm(formula = race_position ~ team_name + wind_speed + rainfall +
##     humidity + pressure + qual_position + loc_cluster + temp_diff,
##     data = train[c("team_name", "wind_speed", "rainfall", "humidity",
##         "pressure", "qual_position", "loc_cluster", "temp_diff",
##         "race_position")])
##
## Residuals:
##      Min      1Q   Median      3Q      Max 
## -13.5659 -2.4499 -0.2054  2.2924 14.4609 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)             5.6973695  0.4846213 11.756 < 2e-16 ***
## team_nameAlphaTauri    -0.4602610  0.1620438 -2.840 0.004510 **  
## team_nameAlpine        -0.3645462  0.1315489 -2.771 0.005590 **  
## team_nameAston Martin  -0.7557428  0.1318753 -5.731 1.01e-08 *** 
## team_nameFerrari       -2.4410510  0.1438583 -16.968 < 2e-16 *** 
## team_nameHaas F1 Team -0.2399908  0.1296182 -1.852 0.064107 .    
## team_nameKick Sauber   0.3745841  0.1422188  2.634 0.008448 **  
## team_nameMcLaren        -2.0402580  0.1438182 -14.186 < 2e-16 *** 
## team_nameMercedes      -1.7280713  0.1396832 -12.371 < 2e-16 *** 
## team_nameRacing Bulls  0.1275074  0.1717690  0.742 0.457902  
## team_nameRB            0.1649315  0.1653202  0.998 0.318460  
## team_nameRed Bull Racing -2.3387322 0.1403298 -16.666 < 2e-16 *** 
## team_nameWilliams       -0.1639976  0.1315355 -1.247 0.212485  
## wind_speed              0.0113492  0.0301699  0.376 0.706789  
## rainfall                0.2097703  0.2713171  0.773 0.439438  
## humidity                -0.0023706  0.0018234 -1.300 0.193573  
## pressure                0.0018329  0.0004721  3.882 0.000104 *** 
## qual_position           0.4557103  0.0052902  86.142 < 2e-16 *** 
## loc_cluster2            -0.2660195  0.0592456 -4.490 7.15e-06 *** 
## temp_diff               -0.0186156  0.0045725 -4.071 4.69e-05 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.538 on 23574 degrees of freedom
## Multiple R-squared:  0.4347, Adjusted R-squared:  0.4342 
## F-statistic:  954 on 19 and 23574 DF,  p-value: < 2.2e-16

```

```

#without qualifying results
lm_race <- lm(race_position ~ ., data = train[c("team_name","wind_speed","rainfall","humidity","pressure","loc_cluster","temp_diff","race_position")])
summary(lm_race)

```

```

## 
## Call:
## lm(formula = race_position ~ ., data = train[c("team_name", "wind_speed",
##   "rainfall", "humidity", "pressure", "loc_cluster", "temp_diff",
##   "race_position")])
##
## Residuals:
##    Min      1Q  Median      3Q     Max
## -12.3283 -3.0199 -0.0893  2.8667 12.5390
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)                12.5829315  0.5480612 22.959 < 2e-16 ***
## team_nameAlphaTauri       -0.3947155  0.1857990 -2.124 0.033645 *  
## team_nameAlpine            -1.0708273  0.1505420 -7.113 1.17e-12 ***
## team_nameAston Martin     -2.0193805  0.1502711 -13.438 < 2e-16 ***
## team_nameFerrari           -6.1347353  0.1574517 -38.963 < 2e-16 ***
## team_nameHaas F1 Team     -0.5266671  0.1485725 -3.545 0.000394 *** 
## team_nameKick Sauber      0.7611597  0.1629883  4.670 3.03e-06 *** 
## team_nameMcLaren           -5.6996766  0.1575452 -36.178 < 2e-16 *** 
## team_nameMercedes          -5.3218814  0.1528517 -34.817 < 2e-16 *** 
## team_nameRacing Bulls     -1.6070348  0.1955941 -8.216 < 2e-16 *** 
## team_nameRB                 -0.7013402  0.1892067 -3.707 0.000210 *** 
## team_nameRed Bull Racing  -5.0293234  0.1568673 -32.061 < 2e-16 *** 
## team_nameWilliams          -0.4610096  0.1507681 -3.058 0.002233 **  
## wind_speed                  0.0098093  0.0345931  0.284 0.776748  
## rainfall                    0.6914275  0.3110289  2.223 0.026223 *  
## humidity                   -0.0007400  0.0020906 -0.354 0.723365  
## pressure                    0.0016104  0.0005413  2.975 0.002934 **  
## loc_cluster2                -0.4113674  0.0679040 -6.058 1.40e-09 *** 
## temp_diff                   -0.0209022  0.0052428 -3.987 6.72e-05 *** 
## ---                        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.056 on 23575 degrees of freedom
## Multiple R-squared:  0.2567, Adjusted R-squared:  0.2562 
## F-statistic: 452.4 on 18 and 23575 DF,  p-value: < 2.2e-16

```

```

lm_race_step <- step(lm_race, scope=list(lower=lm(train$race_position~0), upper=lm_race),
                      direction = "forward")

```

```

## Start: AIC=66093.87
## race_position ~ team_name + wind_speed + rainfall + humidity +
##   pressure + loc_cluster + temp_diff

```

```

summary(lm_race_step)

```

```

## 
## Call:
## lm(formula = race_position ~ team_name + wind_speed + rainfall +
##     humidity + pressure + loc_cluster + temp_diff, data = train[c("team_name",
##     "wind_speed", "rainfall", "humidity", "pressure", "loc_cluster",
##     "temp_diff", "race_position")])
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -12.3283 -3.0199 -0.0893  2.8667 12.5390 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)           12.5829315  0.5480612  22.959 < 2e-16 ***
## team_nameAlphaTauri -0.3947155  0.1857990 -2.124 0.033645 *  
## team_nameAlpine      -1.0708273  0.1505420 -7.113 1.17e-12 ***
## team_nameAston Martin -2.0193805  0.1502711 -13.438 < 2e-16 *** 
## team_nameFerrari     -6.1347353  0.1574517 -38.963 < 2e-16 *** 
## team_nameHaas F1 Team -0.5266671  0.1485725 -3.545 0.000394 *** 
## team_nameKick Sauber   0.7611597  0.1629883  4.670 3.03e-06 *** 
## team_nameMcLaren      -5.6996766  0.1575452 -36.178 < 2e-16 *** 
## team_nameMercedes     -5.3218814  0.1528517 -34.817 < 2e-16 *** 
## team_nameRacing Bulls  -1.6070348  0.1955941 -8.216 < 2e-16 *** 
## team_nameRB            -0.7013402  0.1892067 -3.707 0.000210 *** 
## team_nameRed Bull Racing -5.0293234  0.1568673 -32.061 < 2e-16 *** 
## team_nameWilliams     -0.4610096  0.1507681 -3.058 0.002233 ** 
## wind_speed             0.0098093  0.0345931  0.284 0.776748    
## rainfall               0.6914275  0.3110289  2.223 0.026223 *  
## humidity              -0.0007400  0.0020906 -0.354 0.723365    
## pressure               0.0016104  0.0005413  2.975 0.002934 ** 
## loc_cluster2          -0.4113674  0.0679040 -6.058 1.40e-09 *** 
## temp_diff              -0.0209022  0.0052428 -3.987 6.72e-05 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 4.056 on 23575 degrees of freedom
## Multiple R-squared:  0.2567, Adjusted R-squared:  0.2562 
## F-statistic: 452.4 on 18 and 23575 DF,  p-value: < 2.2e-16

```

```
# Error
train$pred_wqual <- predict(lm_race_step_wqual,train[c("team_name","wind_speed","rainfall",
"humidity","pressure","qual_position","loc_cluster","temp_diff")])
train$pred <- predict(lm_race_step,train[c("team_name","wind_speed","rainfall","humidity",
"pressure","loc_cluster","temp_diff")])

test$pred_wqual <- predict(lm_race_step_wqual,test[c("team_name","wind_speed","rainfall",
"humidity","pressure","qual_position","loc_cluster","temp_diff")])
test$pred <- predict(lm_race_step,test[c("team_name","wind_speed","rainfall","humidity",
"pressure","loc_cluster","temp_diff")])

#Train
accuracy(train$pred_wqual,train$race_position)
```

```
##               ME      RMSE      MAE      MPE      MAPE
## Test set 3.22807e-13 3.536062 2.803745 -21.15566 39.6356
```

```
accuracy(train$pred,train$race_position)
```

```
##               ME      RMSE      MAE      MPE      MAPE
## Test set 1.003642e-12 4.05457 3.318377 -28.79826 49.99304
```

```
#Test
accuracy(test$pred_wqual,test$race_position)
```

```
##               ME      RMSE      MAE      MPE      MAPE
## Test set -0.00467727 3.5437 2.816788 -20.30624 38.85112
```


```
accuracy(test$pred,test$race_position)
```

```
##               ME      RMSE      MAE      MPE      MAPE
## Test set 0.01578868 4.062522 3.341963 -27.29216 48.74729
```

```
#Pretty Good MAE
```

```
# -----
# Decision Tree
# -----
#With qual
reg_tree_wqual <- rpart(race_position ~ ., data = train[c("team_name","wind_speed","rainfall",
"humidity","pressure","qual_position","loc_cluster","temp_diff","race_position")])

plot(as.party(reg_tree_wqual), type="extended")
```




```
printcp(reg_tree_wqual)
```

```
##
## Regression tree:
## rpart(formula = race_position ~ ., data = train[c("team_name",
##   "wind_speed", "rainfall", "humidity", "pressure", "qual_position",
##   "loc_cluster", "temp_diff", "race_position")])
##
## Variables actually used in tree construction:
## [1] qual_position
##
## Root node error: 521840/23594 = 22.117
##
## n= 23594
##
##          CP nsplit rel error  xerror      xstd
## 1 0.321565     0  1.00000 1.00007 0.0069094
## 2 0.035901     1  0.67844 0.67858 0.0059674
## 3 0.031257     2  0.64253 0.64571 0.0060100
## 4 0.011380     3  0.61128 0.61644 0.0059369
## 5 0.010000     4  0.59990 0.60561 0.0058975
```

#Somewhat useless as we see that qualifying position is fairly useless

#No qual

```
reg_tree <- rpart(race_position ~ ., data = train[c("team_name","wind_speed","rainfall","humidity","pressure","loc_cluster","temp_diff","race_position")))  
plot(as.party(reg_tree), type="extended")
```



```
printcp(reg_tree)
```

```
##  
## Regression tree:  
## rpart(formula = race_position ~ ., data = train[c("team_name",  
##     "wind_speed", "rainfall", "humidity", "pressure", "loc_cluster",  
##     "temp_diff", "race_position")])  
##  
## Variables actually used in tree construction:  
## [1] team_name  
##  
## Root node error: 521840/23594 = 22.117  
##  
## n= 23594  
##  
##      CP nsplit rel_error xerror      xstd  
## 1 0.233892      0  1.00000 1.00004 0.0069092  
## 2 0.012264      1  0.76611 0.76618 0.0062343  
## 3 0.010000      2  0.75384 0.75394 0.0061682
```

```

#This time only split on team name

# -----
# Forecast Data for 2026
# -----
forecast <- function(){
cities <- c(
  "Melbourne", "Shanghai", "Suzuka", "Sakhir", "Jeddah", "Miami",
  "Montreal", "Monaco", "Barcelona", "Spielberg", "Silverstone",
  "Spa-Francorchamps", "Budapest", "Zandvoort", "Monza", "Madrid",
  "Baku", "Singapore", "Austin", "Mexico City", "São Paulo",
  "Las Vegas", "Lusail", "Yas Marina"
)
start_dates_str <- c(
  "2026-03-06", "2026-03-13", "2026-03-27", "2026-04-10",
  "2026-04-17", "2026-05-01", "2026-05-22", "2026-06-05",
  "2026-06-12", "2026-06-26", "2026-07-03", "2026-07-17",
  "2026-07-24", "2026-08-21", "2026-09-04", "2026-09-11",
  "2026-09-24", "2026-10-09", "2026-10-23",
  "2026-10-30", "2026-11-06", "2026-11-19",
  "2026-11-27", "2026-12-04"
)
latitude <- c(
  -37.81, 31.23, 34.84, 26.06, 21.53, 25.76,
  45.50, 43.74, 41.39, 47.21, 52.09,
  50.44, 47.50, 52.38, 45.58, 40.42,
  40.41, 1.35, 30.27, 19.43, -23.70,
  36.17, 25.42, 24.48
)
longitude <- c(
  144.96, 121.47, 136.54, 50.53, 39.16, -80.19,
  -73.57, 7.42, 2.17, 14.80, -1.03,
  5.97, 19.04, 4.53, 9.27, -3.70,
  49.87, 103.82, -97.74, -99.13, -46.71,
  -115.14, 51.50, 54.61
)

race_dates_converted <- as.Date(start_dates_str)

races_2026 <- data.frame(
  lag_1 = race_dates_converted %m-% years(1),
  lag_2 = race_dates_converted %m-% years(2),
  lag_3 = race_dates_converted %m-% years(3),
  lag_4 = race_dates_converted %m-% years(4),
  lag_5 = race_dates_converted %m-% years(5),
  Latitude = latitude,
  Longitude = longitude,
  stringsAsFactors = FALSE
)

```

```

# -----
# Import Future Weather Data
# -----
lag_weather <- rbind(
  select(races_2026, begin_date = lag_1, Latitude, Longitude) %>%
  pmap_df(.f =
    function(begin_date, Latitude, Longitude){
      url <- paste0("https://archive-api.open-meteo.com/v1/archive?latitude=", Latitude, "&longitude=", Longitude, "&start_date=", begin_date, "&end_date=", begin_date, "&hourly=temperature_2m,relative_humidity_2m,precipitation,soil_temperature_0_to_7cm,wind_speed_10m,surface_pressure")
      df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
      hourly_df <- df$hourly
      hourly_df$longitude <- Longitude
      hourly_df$latitude <- Latitude
      return(hourly_df)
    }),
  select(races_2026, begin_date = lag_2, Latitude, Longitude) %>%
  pmap_df(.f =
    function(begin_date, Latitude, Longitude){
      url <- paste0("https://archive-api.open-meteo.com/v1/archive?latitude=", Latitude, "&longitude=", Longitude, "&start_date=", begin_date, "&end_date=", begin_date, "&hourly=temperature_2m,relative_humidity_2m,precipitation,soil_temperature_0_to_7cm,wind_speed_10m,surface_pressure")
      df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
      hourly_df <- df$hourly
      hourly_df$longitude <- Longitude
      hourly_df$latitude <- Latitude
      return(hourly_df)
    }),
  select(races_2026, begin_date = lag_3, Latitude, Longitude) %>%
  pmap_df(.f =
    function(begin_date, Latitude, Longitude){
      url <- paste0("https://archive-api.open-meteo.com/v1/archive?latitude=", Latitude, "&longitude=", Longitude, "&start_date=", begin_date, "&end_date=", begin_date, "&hourly=temperature_2m,relative_humidity_2m,precipitation,soil_temperature_0_to_7cm,wind_speed_10m,surface_pressure")
      df <- fromJSON(content(GET(url), "text", encoding = "UTF-8"))
      hourly_df <- df$hourly
      hourly_df$longitude <- Longitude
      hourly_df$latitude <- Latitude
      return(hourly_df)
    }),
  select(races_2026, begin_date = lag_4, Latitude, Longitude) %>%
  pmap_df(.f =
    function(begin_date, Latitude, Longitude){

```

```

url <- paste0("https://archive-api.open-meteo.com/
v1/archive?latitude=", Latitude, "&longitude=", Longitude, "&start_date=", begin_date, "&end_d
ate=", begin_date, "&hourly=temperature_2m,relative_humidity_2m,precipitation,soil_tempera
ture_0_to_7cm,wind_speed_10m,surface_pressure")
df <- fromJSON(content(GET(url), "text", encoding
= "UTF-8"))

hourly_df <- df$hourly
hourly_df$longitude <- Longitude
hourly_df$latitude <- Latitude
return(hourly_df)
}),

select(races_2026,begin_date = lag_5, Latitude, Longitude) %>%
  pmap_df(.f=
    function(begin_date, Latitude, Longitude){
      url <- paste0("https://archive-api.open-meteo.com/
v1/archive?latitude=", Latitude, "&longitude=", Longitude, "&start_date=", begin_date, "&end_d
ate=", begin_date, "&hourly=temperature_2m,relative_humidity_2m,precipitation,soil_tempera
ture_0_to_7cm,wind_speed_10m,surface_pressure")
      df <- fromJSON(content(GET(url), "text", encoding
= "UTF-8"))

hourly_df <- df$hourly
hourly_df$longitude <- Longitude
hourly_df$latitude <- Latitude
return(hourly_df)
})))

temp <- lag_weather
lag_weather <- temp

lag_weather <- lag_weather[-1] %>%
  group_by(longitude, latitude) %>%
  summarise(temp_diff = mean(soil_temperature_0_to_7cm) - mean(temperature
_2m),
            humidity = mean(relative_humidity_2m),
            rainfall = mean(precipitation),
            wind_speed = mean(wind_speed_10m),
            pressure = mean(surface_pressure)) %>%
  filter(latitude == max(latitude))

# -----
# Full 2026 Schedule
# -----
schedule_2026 <- data.frame(location = cities, latitude, longitude) %>%
  left_join(lag_weather, join_by(longitude, latitude))
schedule_2026 <- schedule_2026[c(-2,-3)]

```

```
teams <- unique(analysis_set$team_name)

forecast_2026 <- data.frame()
for (t in teams) {
  temp <- schedule_2026
  temp$team_name <- t
  forecast_2026 <- rbind(forecast_2026,temp)
}
return(forecast_2026)
}
forecast_2026 <- forecast()
```

```
## `summarise()` has grouped output by 'longitude'. You can override using the
## `.`groups` argument.
```

```

#write.csv(forecast_2026,"2026_projected_data")
#forecast_2026 <- read.csv("2026_projected_data")
#Shiv

##
##Add Location Cluster
##
forecast_2026 <- left_join(forecast_2026,cluster_assignments,by="location") %>%
select("team_name","wind_speed","rainfall","humidity","pressure",loc_cluster="Track_Cluster","temp_diff","location")

# ##
#Note Montreal Madrid Singapore are new so they do not have a location cluster
#Will add these in
# ##
centroids <- final_clusters$centers
missing_clust <- forecast_2026[is.na(forecast_2026$loc_cluster),] %>%
  select("wind_speed",
         "rainfall",
         "humidity",
         "pressure",
         "temp_diff",
         "location") %>%
  distinct(location, .keep_all = TRUE)

missing_clust$dist1 <- sqrt((missing_clust$wind_speed - centroids[1,1])^2 +
                           (missing_clust$rainfall - centroids[1,2])^2 +
                           (missing_clust$humidity - centroids[1,5])^2 +
                           (missing_clust$pressure - centroids[1,6])^2 +
                           (missing_clust$temp_diff - centroids[1,7])^2)
missing_clust$dist2 <- sqrt((missing_clust$wind_speed - centroids[2,1])^2 +
                           (missing_clust$rainfall - centroids[2,2])^2 +
                           (missing_clust$humidity - centroids[2,5])^2 +
                           (missing_clust$pressure - centroids[2,6])^2 +
                           (missing_clust$temp_diff - centroids[2,7])^2)
missing_clust$loc_cluster <- if_else(missing_clust$dist1 <= missing_clust$dist2,1,2)

missing_clust <- select(missing_clust,"loc_cluster","location")

forecast_2026 <- left_join(forecast_2026,missing_clust,by="location")
forecast_2026$loc_cluster <- if_else(!is.na(forecast_2026$loc_cluster.x),forecast_2026$loc_cluster.x,forecast_2026$loc_cluster.y)
forecast_2026 <- select(forecast_2026,-"loc_cluster.x",-"loc_cluster.y")

# ##
#Make forecast
# ##
forecast_2026$team_name <- as.factor(forecast_2026$team_name)
forecast_2026$loc_cluster <- as.factor(forecast_2026$loc_cluster)

forecast_2026$predicted_position <- predict(lm_race_step,forecast_2026)

```

```
standing_table <- group_by(forecast_2026,team_name) %>%
  summarise(pos_sum = sum(predicted_position))

standing_table$predicted_rank <- rank(standing_table$pos_sum)
```